




# Learning Dirichlet Priors for Affordance Aware Planning

David Abel, Gabriel Barth-Maron, James MacGlashan, Stefanie Tellex  
Dept. of Computer Science, Brown University




## Goal

**Previous Work:** Provided an MDP with knowledge in order to solve extremely complex, previously unsolved tasks.

**Proposal:** Learn this knowledge to remove dependence on expert.

## Background

**Affordances:** Direct agent toward relevant action possibilities.



“What [the environment] offers [an] animal, what [the environment] provides or furnishes, either for good or ill”

- J.J. Gibson, 1977

Formalism:

$$\Delta = \langle p, g \rangle \mapsto \mathcal{A}'$$

$p$  = predicate on states

$g$  = lifted goal description

$\mathcal{A}'$  = subset of OO-MDP Actions

**Domain:** Minecraft




$\approx$  Turing Complete Legos

## Affordance Example

$$\Delta_1 = \langle \text{nearPlane}, \text{atLoc} \rangle \mapsto \{\text{move}\}$$

$$\Delta_2 = \langle \text{nearTrench}, \text{atLoc} \rangle \mapsto \{\text{place}\}$$

If  $\Delta$ 's predicate is true and  $\Delta$ 's goal type matches the current goal, use  $\Delta$ 's actions




## Learning

**Goal:** For a given state, for each affordance, learn which actions are most relevant:

$$\Pr(\mathcal{A}^* \mid s, \Delta_1 \dots \Delta_K)$$

**Graphical Model:**



Where:

$$\Pr(\lambda \mid \alpha) = \text{DirMult}(\alpha)$$

$$\Pr(n \mid \beta) = \text{Dir}(\beta)$$

$s$  = OO-MDP State

$\Delta$  = Affordance

$\alpha$  = Action Counts

$\beta$  = Action Set Size Counts

$\lambda$  = Distribution on Actions

$n$  = Distribution on Action Set Size

$\mathcal{A}'$  = One Affordance's Action Set

$$\mathcal{A}^* = \bigcup_{i=1}^K \mathcal{A}'_i$$


## Learning Example

1) For each activated affordance, count:

$\alpha$  = number of worlds in which each action was used

$\beta$  = number of unique actions used in each world

→ = move  
→ = place



$$\Delta_i.\alpha \leftarrow \{\text{moveRight++}, \text{moveForward++}, \text{placeRight++}\}$$

$$\Delta_i.\beta \leftarrow \{3++\}$$

2) When solving the MDP on a new state space, in each state  $s$ :

$$\mathcal{A}^* = \bigcup_{i=1}^K (\Delta_i.\text{getActions}(s))$$

3) Where

$\Delta_i.\text{getActions}(s)$ :

$\lambda \leftarrow \text{DirMult}(\Delta_i.\alpha)$   
 $n \leftarrow \text{Dir}(\Delta_i.\beta)$   
 $\mathcal{A}' \leftarrow_n \lambda$   
return:  $\mathcal{A}'$

## Results

**Avg. # Bellman Updates Per Converged Policy**

