




Learning Dirichlet Priors for Affordance Aware Planning

David Abel, Gabriel Barth-Maron, James MacGlashan, Stefanie Tellex
Dept. of Computer Science, Brown University




Goal

Previous Work: Provided an MDP with knowledge in order to solve extremely complex, previously unsolved tasks.

Proposal: Learn this knowledge to remove dependence on expert.

Background

Affordances: Direct agent toward relevant action possibilities.



“What [the environment] offers [an] animal, what [the environment] provides or furnishes, either for good or ill”

- J.J. Gibson, 1977

Formalism:

$$\Delta = \langle p, g \rangle \mapsto \mathcal{A}'$$

p = predicate on states

g = lifted goal description

\mathcal{A}' = subset of OO-MDP Actions

Domain: Minecraft




\approx Turing Complete Legos

Affordance Example

$$\Delta_1 = \langle \text{nearPlane}, \text{atLoc} \rangle \mapsto \{\text{move}\}$$

$$\Delta_2 = \langle \text{nearTrench}, \text{atLoc} \rangle \mapsto \{\text{place}\}$$

If Δ 's predicate is true and Δ 's goal type matches the current goal, use Δ 's actions




Learning

Goal: For a given state, for each affordance, learn which actions are most relevant:

$$\Pr(\mathcal{A}^* \mid s, \Delta_1 \dots \Delta_K)$$

Graphical Model:



s = OO-MDP State

Δ = Affordance

α = Action Counts

β = Action Set Size Counts

λ = Distribution on Actions

n = Distribution on Action Set Size

\mathcal{A}' = One Affordance's Action Set

$$\mathcal{A}^* = \bigcup_{i=1}^K \mathcal{A}'_i$$

Where:

$$\Pr(\lambda \mid \alpha) = \text{DirMult}(\alpha)$$


$$\Pr(n \mid \beta) = \text{Dir}(\beta)$$

Learning Example

1. For each activated affordance, count:

α = number of worlds in which each action was used

β = number of unique actions used in each world



$$\Delta_i.\alpha \leftarrow \{\text{moveRight++}, \text{moveForward++}, \text{placeRight++}\}$$

$$\Delta_i.\beta \leftarrow \{3++\}$$

2. We have:

$\Delta_i.\text{getActions}(s)$:

```

λ ← DirMult(Δi.α)
n ← Dir(Δi.β)
A' ←n λ
return: A'

```

3. When solving the MDP on a new state space, in each state s :

$$\mathcal{A}^* = \bigcup_{i=1}^K (\Delta_i.\text{getActions}(s))$$

Results

Avg. # Bellman Updates Per Converged Policy

